



National  
Taiwan  
University

**CSIE3310 - Operating Systems**

# **Machine Problem 2**

## **Demand Paging and Swapping**

**Department of Computer Science and Information Engineering**

TA e-mail: [ntuos@googlegroups.com](mailto:ntuos@googlegroups.com)

TA hours: Tue. & Thu. 13:00-14:00 CSIE Building R442 or B04 (Please knock the door)

# 01 / TA Contact

## **林祥瑞 / Jerry Lin**

- d10922013@ntu.edu.tw
- At room B04

## **劉昕祐 / Xin-You Liu**

- r10944004@csie.ntu.edu.tw
- At room R442

Better to ask your questions on homework discussion page on NTUCOOL instead of emails.

# Outline

- PART 01 / Summary
- PART 02 / Launching
- PART 03 / Assignment
- PART 04 / Test Programs
- PART 05 / Submission



PART ONE

# Summary

# 01 / Summary

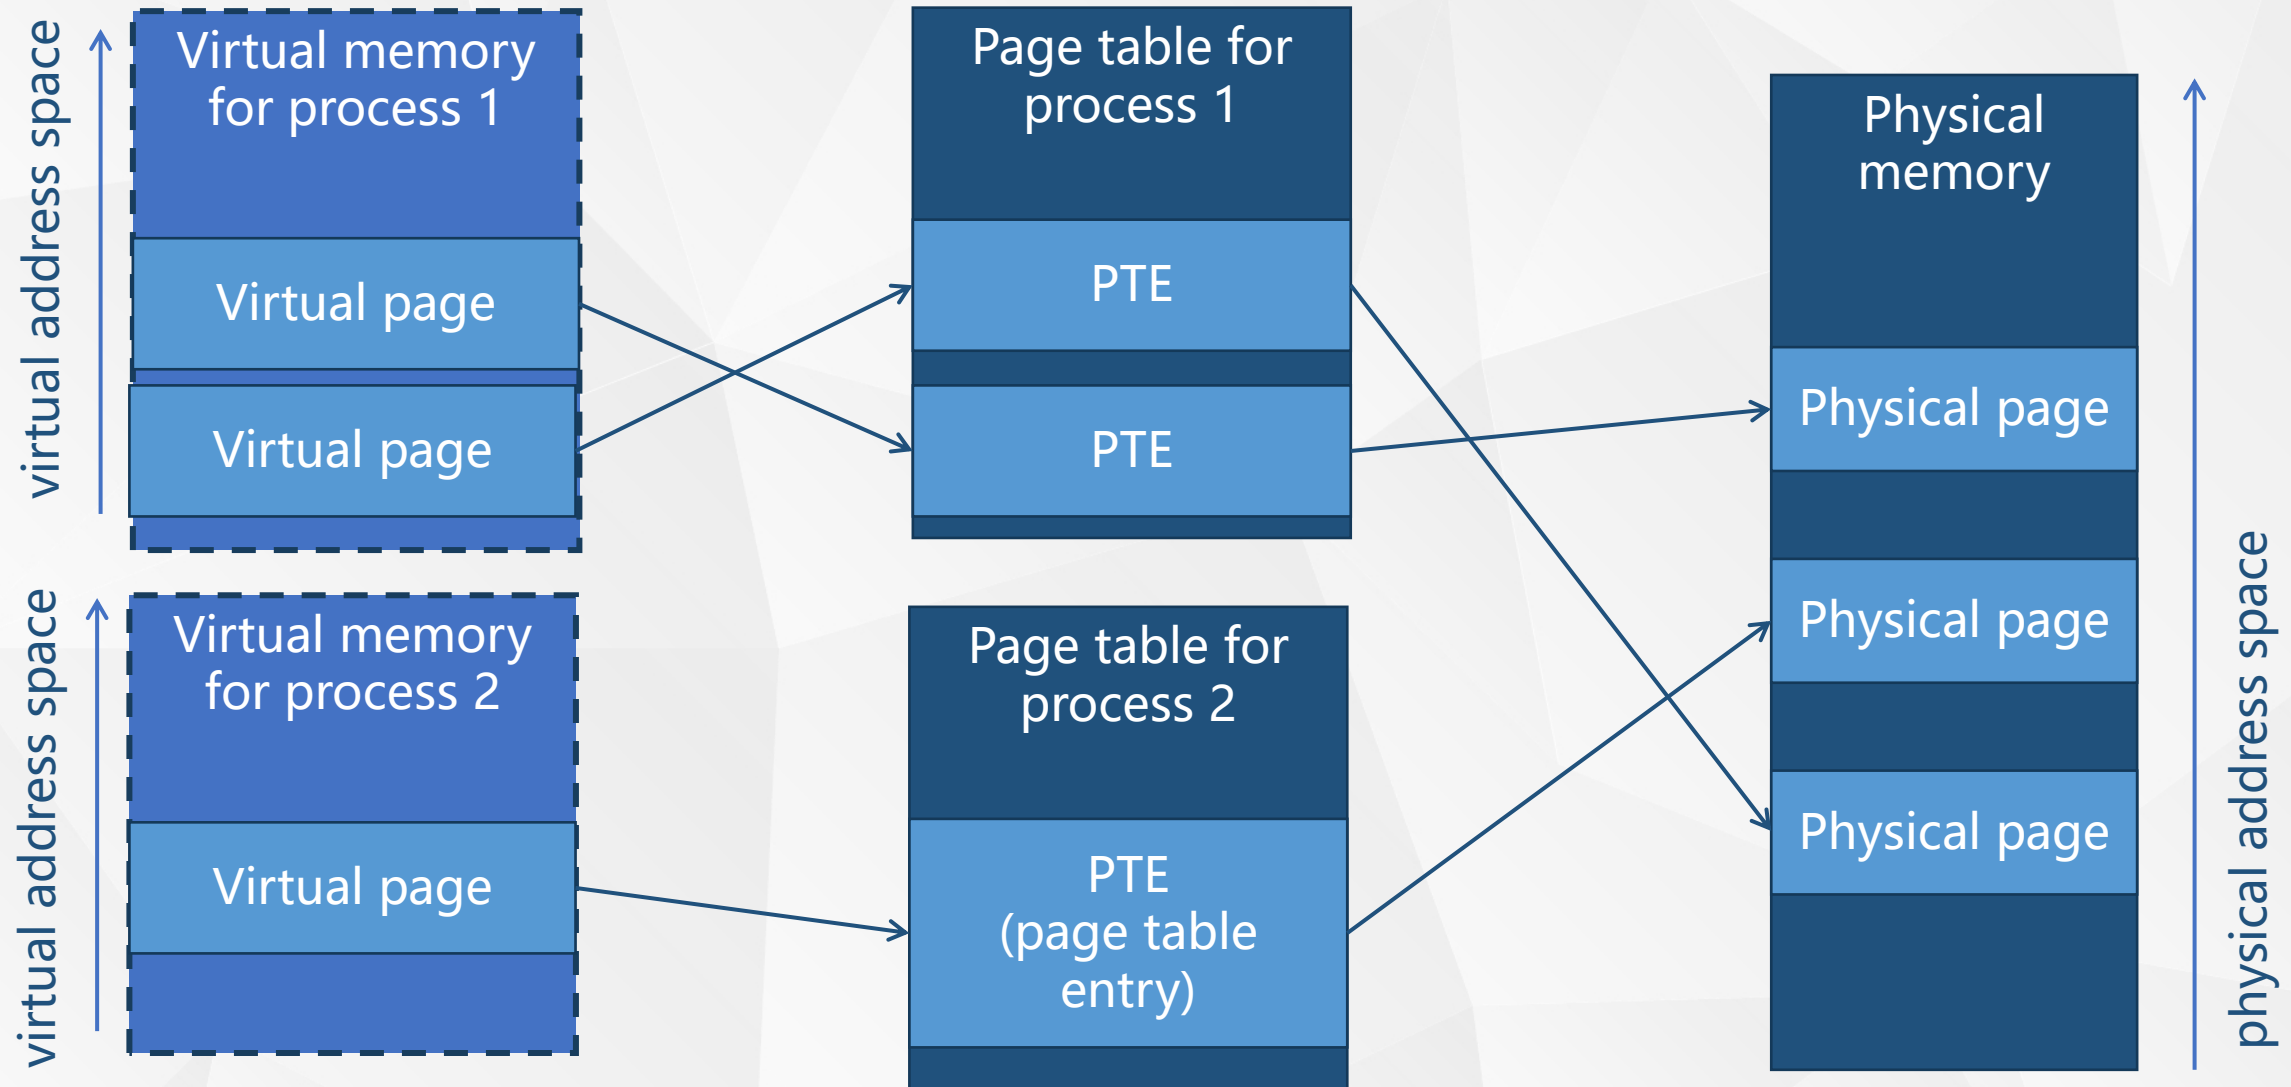
## ❖ Virtual Memory:

- An isolated and abstracted memory space for each process.
- A portion of VM pages are mapped to physical memory through per-process page table.

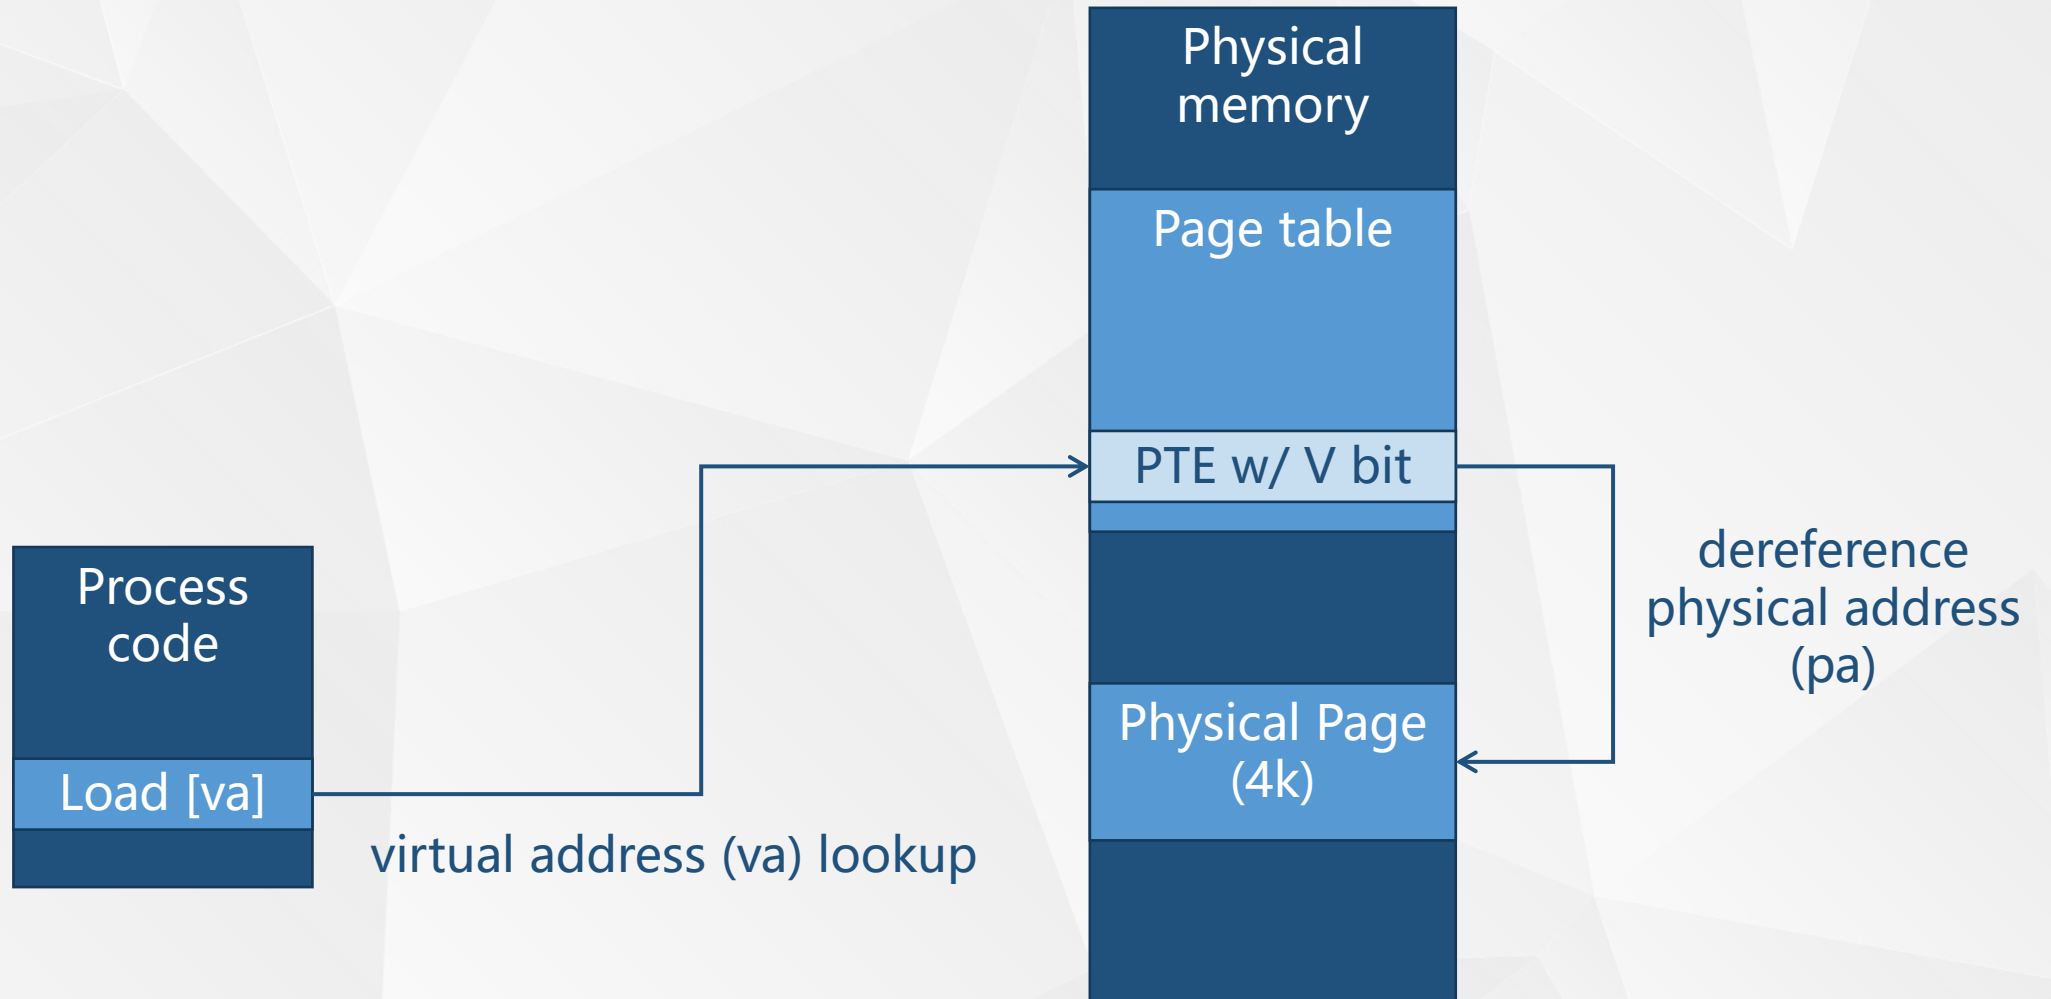
## ❖ Swapping :

- Allow VM pages not only mapped to physical memory pages, can also be mapped to blocks on a disk.
- The OS can swap out “cold” memory pages to disk blocks, or swap in disk blocks to physical memory when needed.

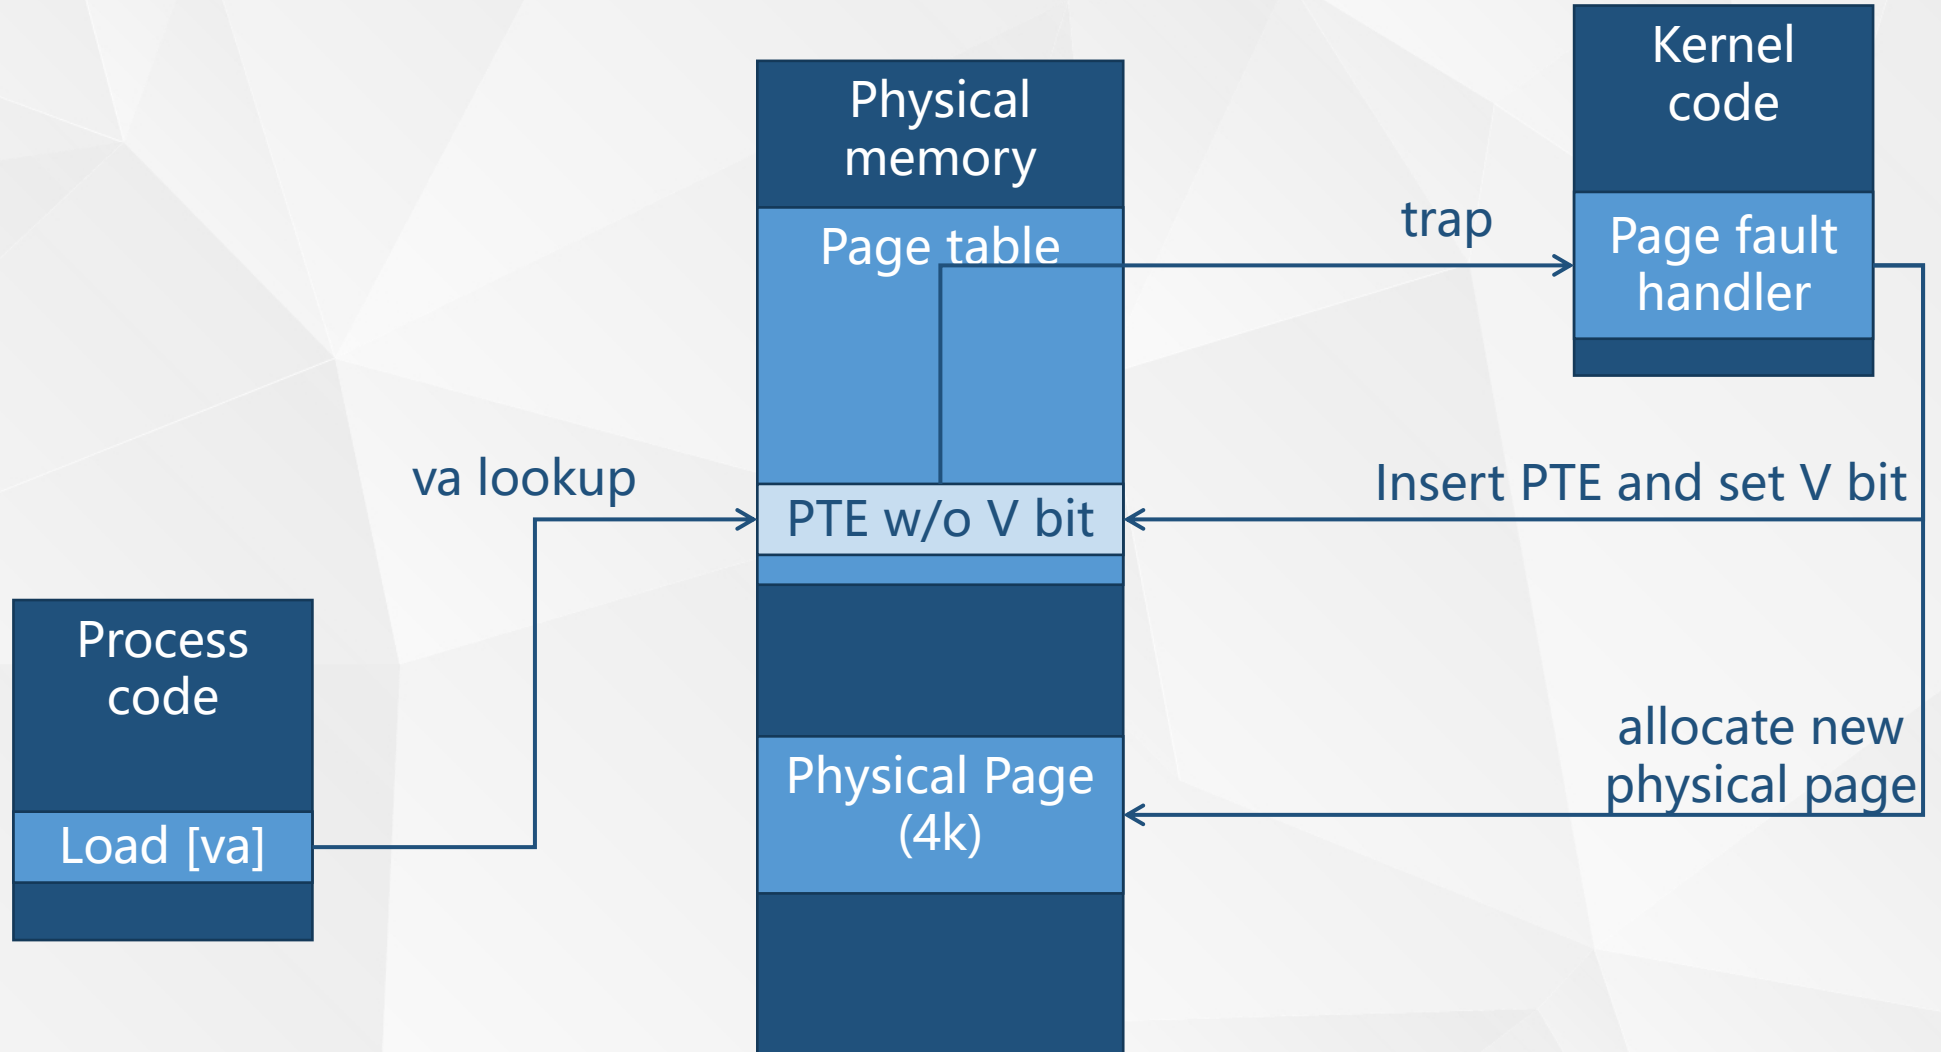
# 01 / Logical View of Virtual Memory



# 01 / Page Hit Workflow

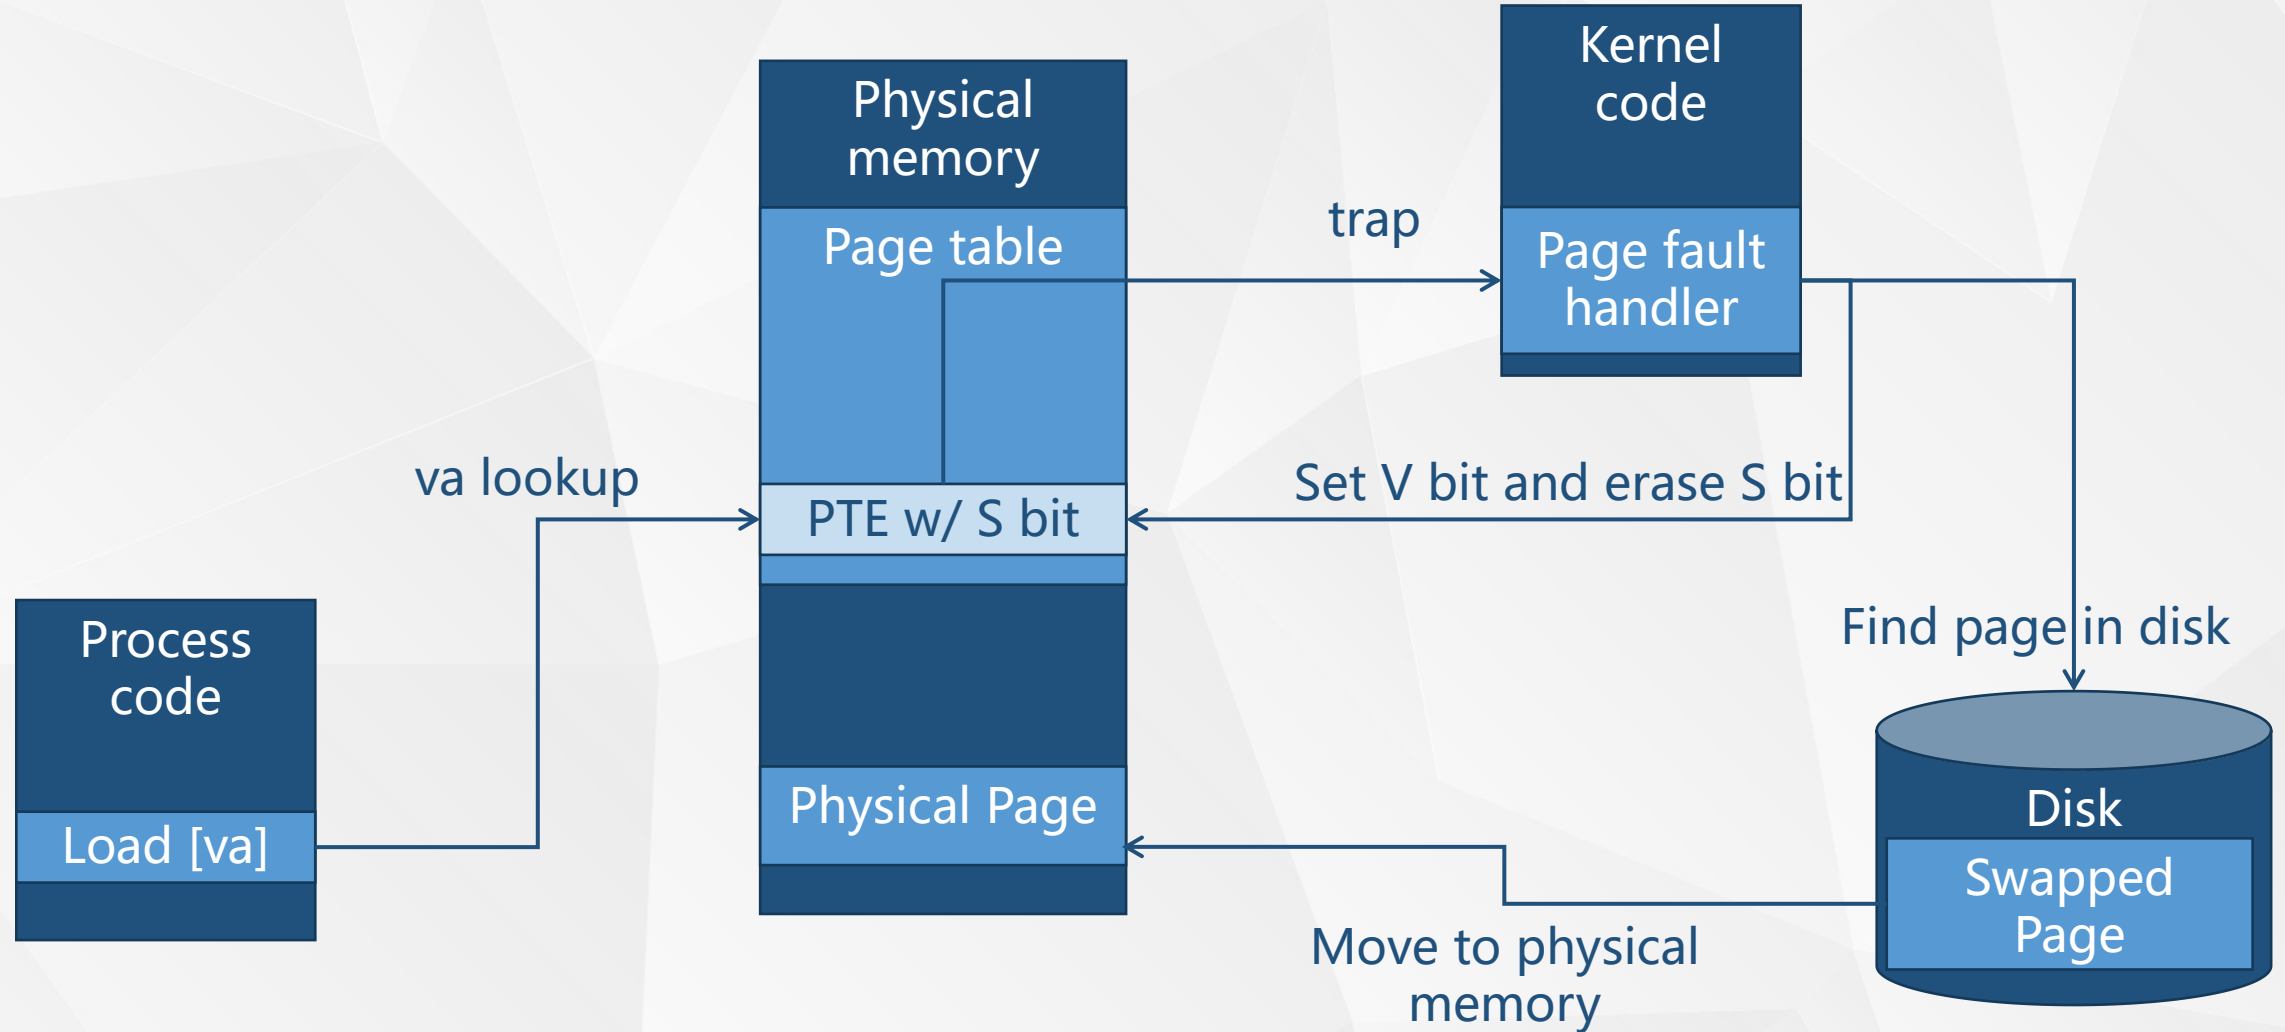


# 01 / Page Miss Workflow

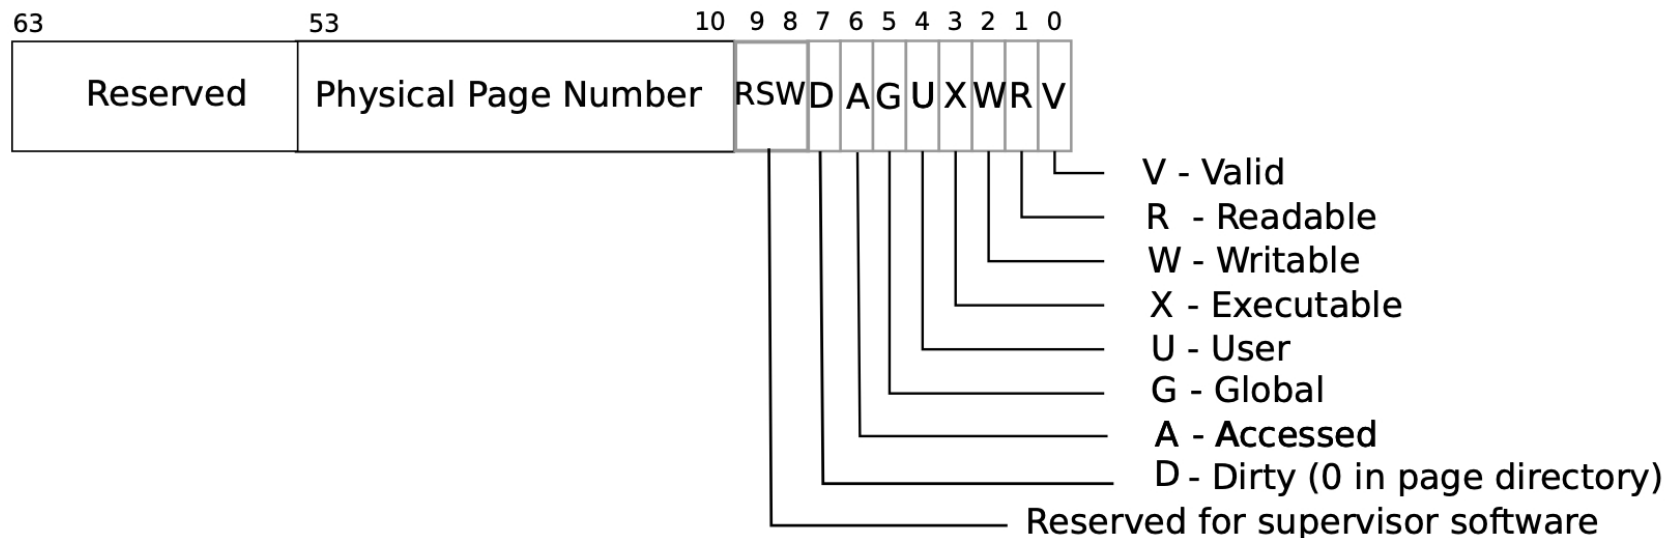
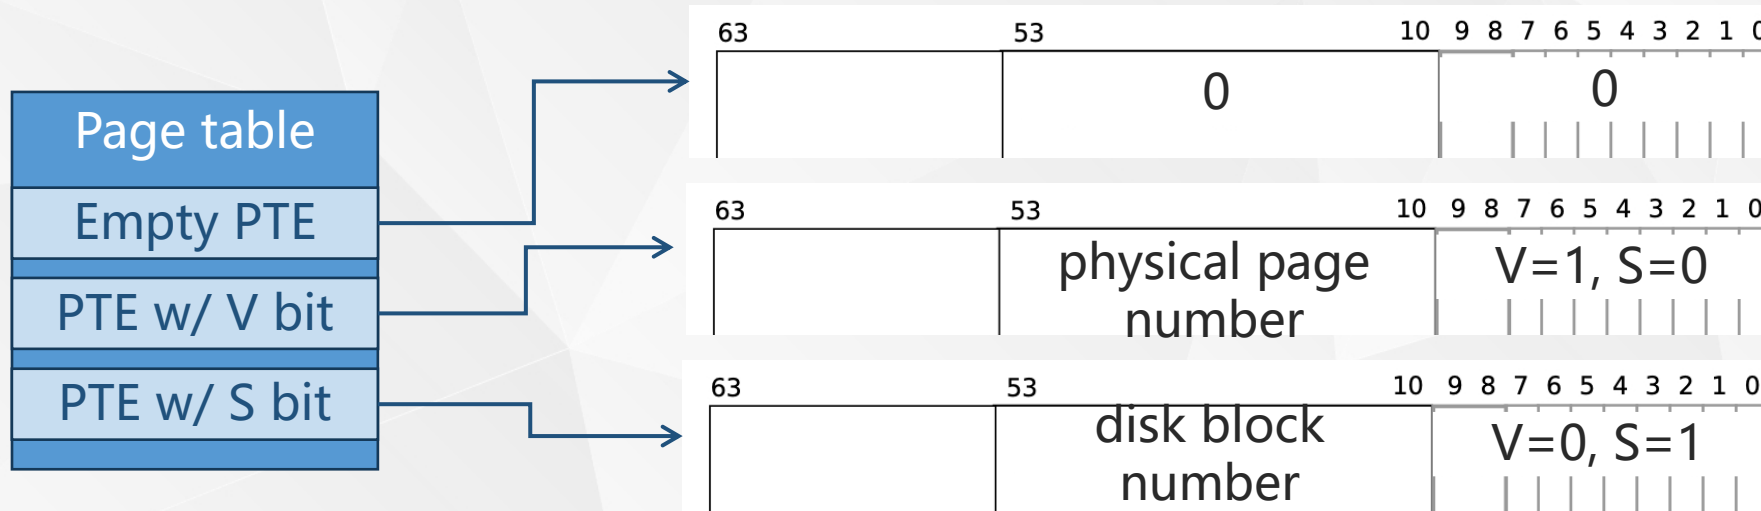




# 01 / Page Miss w/ Swap



# 01 / Page Table Entry Format



# 01 / Summary

## ❖ Machine Problem 2 Goal:

- **Demand Paging :**

The physical memory is allocated on demand only when the virtual pages are accessed.

- **Swapping :**

A technique to store memory pages on a disk.

- With proper memory management, OS can maintain processes with large virtual memory spaces but small physical memory in use.

- MP2 assignment will add “Demand Paging” and “Swapping” to existing page table on xv6.



**02**

**PART TWO**  
**Launching**

# 02 / Launching

## Launching Docker

- **Launching Docker Image of MP2**

1. Download the MP2.zip from NTUCOOL, unzip it, and enter it.

```
$ unzip MP2.zip  
$ cd mp2
```

2. Pull Docker image from Docker Hub

```
$ docker pull ntuos/mp2
```

3. Use docker run to start the process in a container and allocate a TTY for the container process.

```
$ docker run -it -v $(pwd):/home/mp2/xv6 ntuos/mp2
```

4. Check the environment in the Docker container

```
$ cat /etc/os-release
```



**03**

**PART THREE**  
**Assignment**

# 03 / Scoring

- Public program tests (70%)
  - Private program tests (15%)
  - Required report (15%)
  - Bonus report (10%)
- Public program test code is shipped with MP2.zip.
- Private program tests will be disclosed after the deadline.

# 03 / Assignment

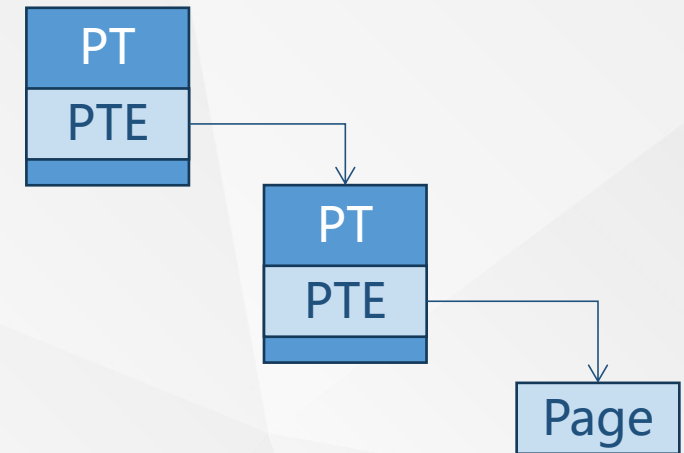
## 1. Print a Page Table (Public Test 5% + Report 5%)

- Traverse the 3-level page table on xv6.
- Print physical addresses and flags, and compute corresponding virtual addresses.

```
xv6 kernel is booting

hart 1 starting
hart 2 starting
init: starting sh
$ mp2_1
page table 0x0000000087f57000
├─ 0: pte=0x0000000087f57000 va=0x0000000000000000 pa=0x0000000087f53000 V
│   └─ 0: pte=0x0000000087f53000 va=0x0000000000000000 pa=0x0000000087f52000 V
│       ├── 0: pte=0x0000000087f52000 va=0x0000000000000000 pa=0x0000000087f54000 V R W X U
│       ├── 1: pte=0x0000000087f52008 va=0x0000000000000100 pa=0x0000000087f51000 V R W X
│       └─ 2: pte=0x0000000087f52010 va=0x0000000000000200 pa=0x0000000087f50000 V R W X U
└─ 255: pte=0x0000000087f577f8 va=0xffffffffc0000000 pa=0x0000000087f56000 V
      └─ 511: pte=0x0000000087f56ff8 va=0xfffffffffe000000 pa=0x0000000087f55000 V
            ├── 510: pte=0x0000000087f55ff0 va=0xfffffffffff000 pa=0x0000000087f65000 V R W
            └─ 511: pte=0x0000000087f55ff8 va=0xfffffffffff000 pa=0x0000000080007000 V R X

$ qemu-system-riscv64: terminating on signal 15 from pid 147755 (make)
```





# 03 / Assignment

## 1. Print a Page Table (Public Test 5% + Report 5%)

Report Part: (5%) Please answer the following questions:

```
xv6 kernel is booting
```

```
hart 1 starting
```

```
hart 2 starting
```

```
init: starting sh
```

```
$ mp2_1
```

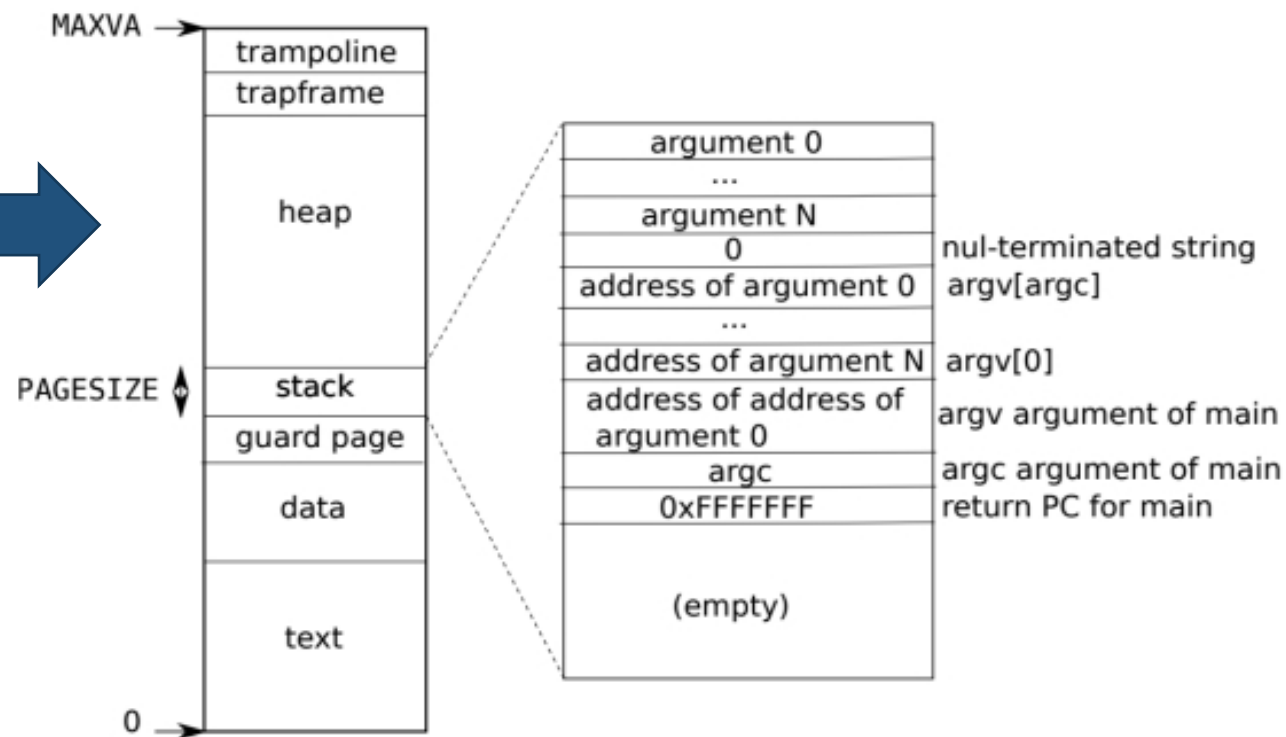
```
page table 0x0000000087f57000
```

```
0: pte=0x0000000087f57000 va=0x0000000000000000 pa=0x0000000087f53000 V
├─ 0: pte=0x0000000087f53000 va=0x0000000000000000 pa=0x0000000087f52000 V
│   ├─ 0: pte=0x0000000087f52000 va=0x0000000000000000 pa=0x0000000087f54000 V R W X U
│   └─ 1: pte=0x0000000087f52008 va=0x0000000000000100 pa=0x0000000087f51000 V R W X
│       └─ 2: pte=0x0000000087f52010 va=0x0000000000000200 pa=0x0000000087f50000 V R W X U
└─ 255: pte=0x0000000087f577f8 va=0xffffffffc0000000 pa=0x0000000087f56000 V
      └─ 511: pte=0x0000000087f56ff8 va=0xfffffffffe000000 pa=0x0000000087f55000 V
            └─ 510: pte=0x0000000087f55ff0 va=0xffffffffffe000 pa=0x0000000087f65000 V R W
                └─ 511: pte=0x0000000087f55ff8 va=0xfffffffffff000 pa=0x0000000080007000 V R X
```

```
$ qemu-system-riscv64: terminating on signal 15 from pid 147755 (make)
```

Explain

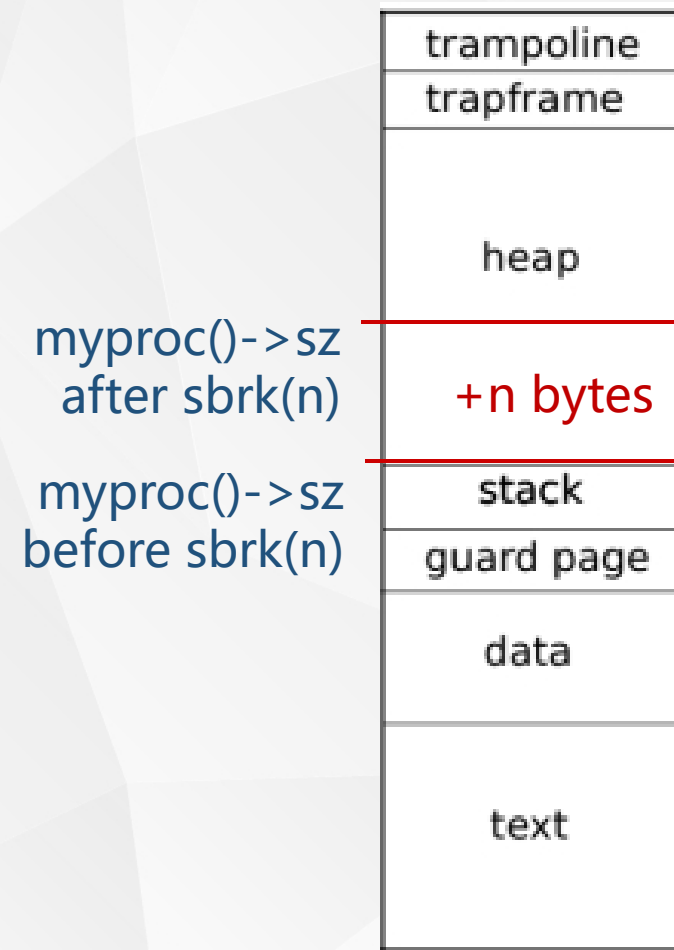
Figure 1: The virtual memory layout



# 03 / Assignment

## 2. Generate a Page Fault (Public Test 20%)

- `sbrk(n)` adds  $n$  bytes to process memory if  $n > 0$ , or decreases by  $|n|$  bytes if  $n < 0$ .
- Change `sbrk(n)` so that it does not allocate.
- Implement page fault handler so that physical pages are allocated only when page fault.



**Original:** Pages are allocated for  $+n$  bytes.

**Changed:** Do not allocate.

# 03 / Assignment

## 3. Demand Paging and Swapping (Public Test 45% + Private Test 15% + Report 10%)

Add the `madvice()` system call with this signature:

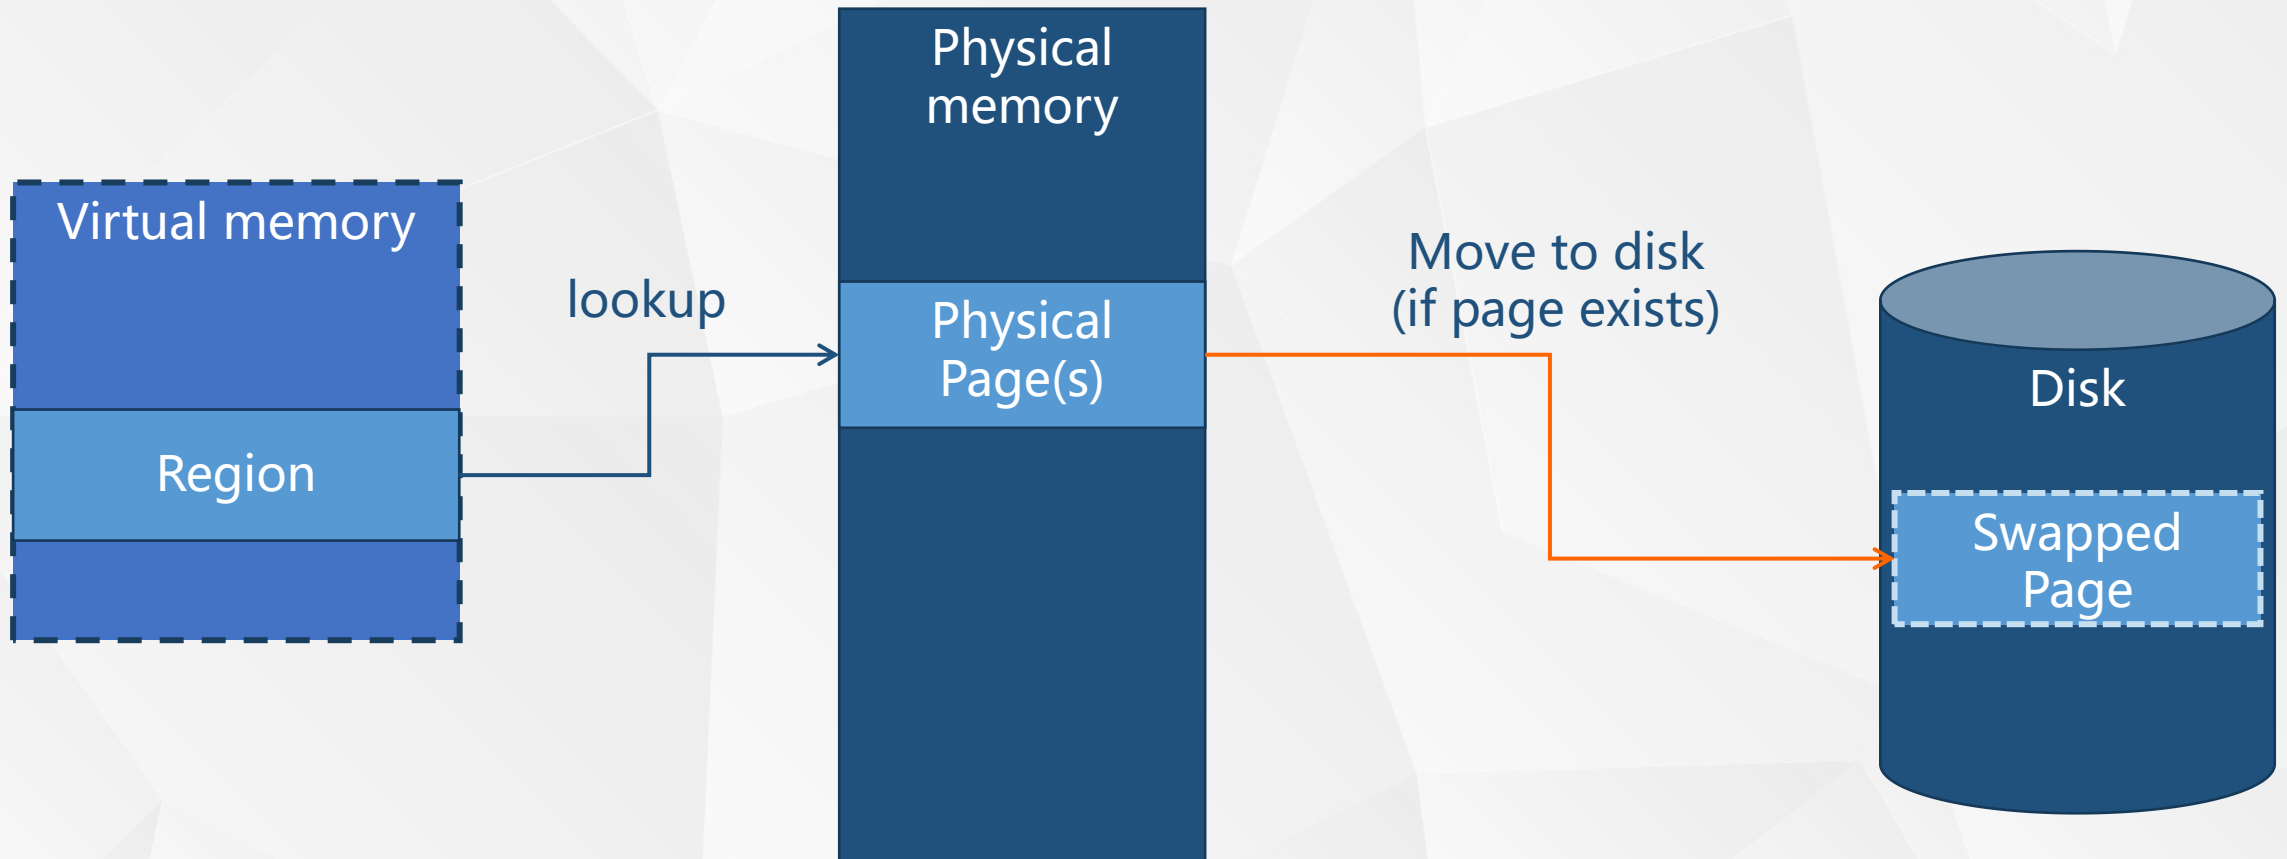
```
int madvice(char *base, int length, int option)
```

```
option :=    MADV_NORMAL,  
             MADV_WILLNEED,  
             MADV_DONTNEED
```

**base** and **length** describes the [base, base + length) virtual memory region in bytes. Return 0 if success, otherwise -1. Invalid range should return -1.

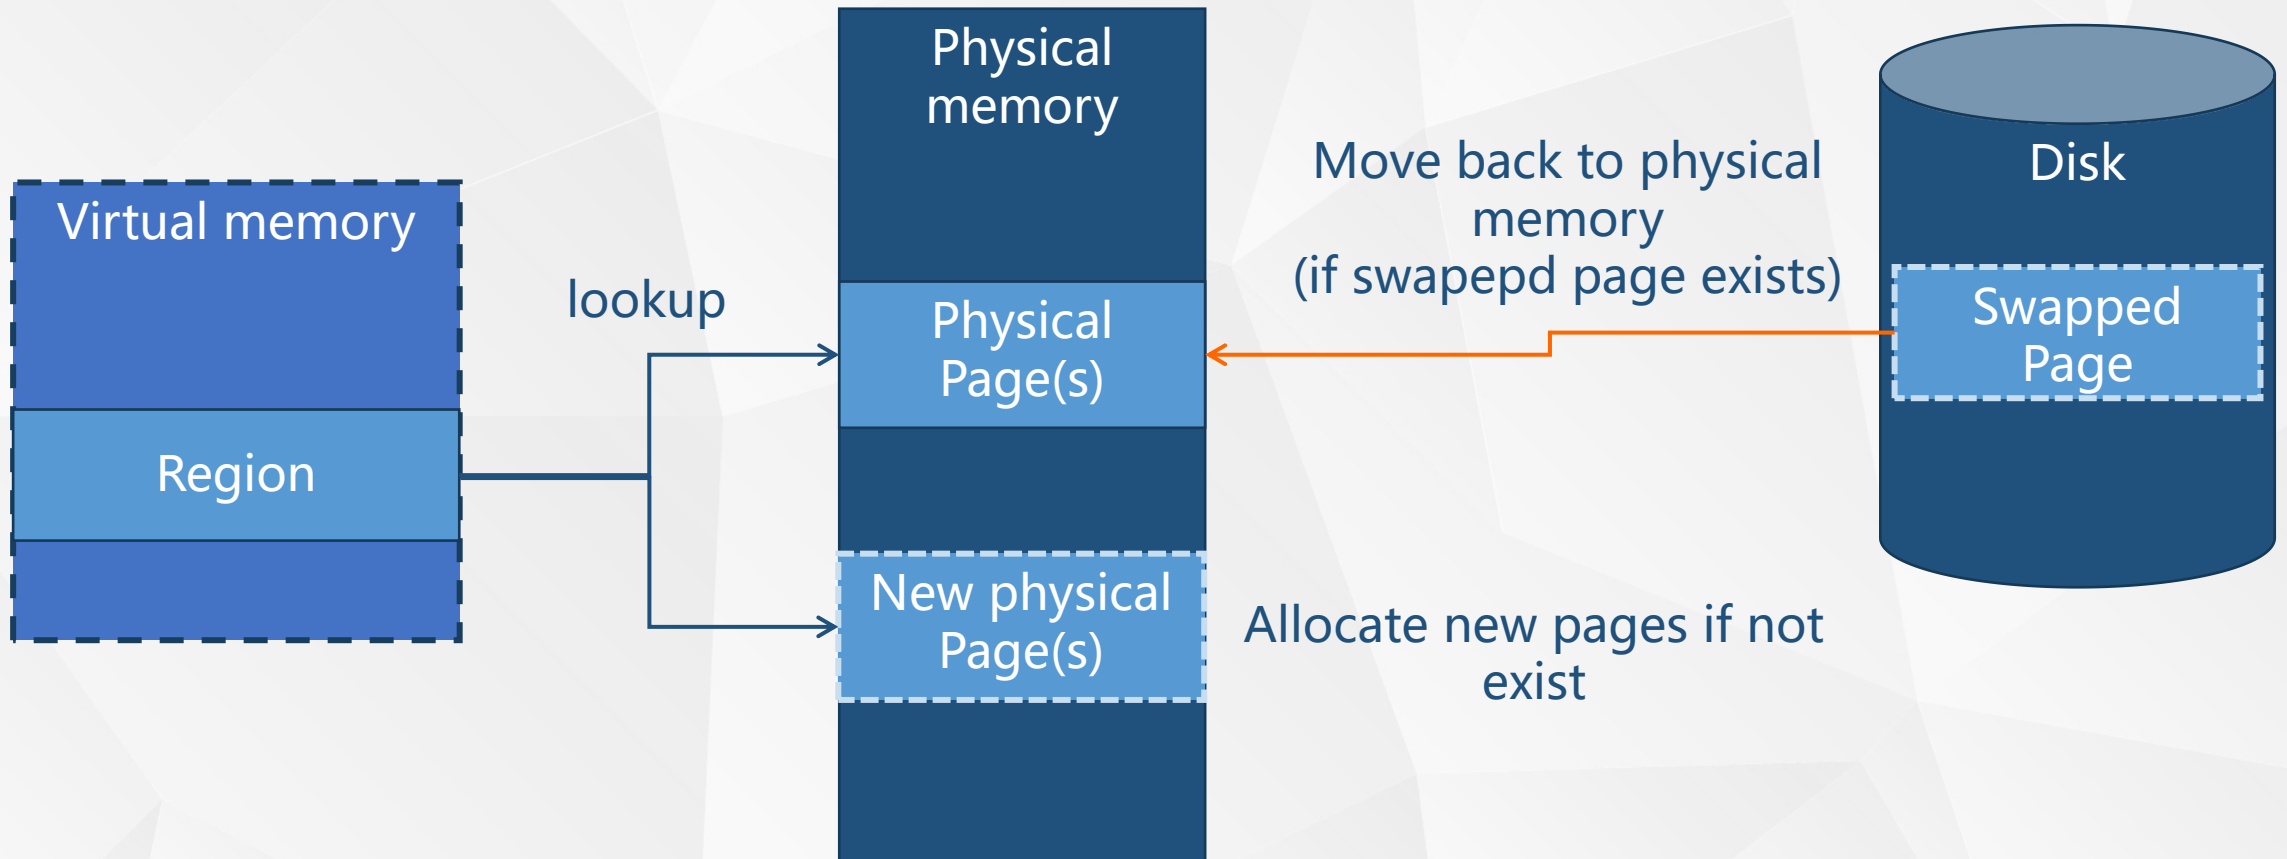
# 03 / Assignment

**MADV\_DONTNEED:** Move affected pages to the disk.



# 03 / Assignment

**MADV\_WILLNEED:** Move affected pages from the disk to physical memory.



# 03 / Assignment

## 3. Demand Paging and Swapping (Public Test 45% + Private Test 15% + Report 10%)

### Program Part I (Public Test 17%)

Test **MADV\_NORMAL** and range checks in `madvise()` + **MADV\_DONTNEED** test

### Program Part II (Public Test 28%)

**MADV\_DONTNEED** + **MADV\_WILLNEED** test

### Program Part III (Private Test 15%)

**MADV\_DONTNEED** + **MADV\_WILLNEED** + page fault on swapped pages

### Report (10%)

# 03 / Assignment

## **4. Bonus Reports (10%)**

**Pros and Cons of Demand Paging (Bonus + 5%):**

**Effective Memory Access Time Analysis (Bonus + 5%)**



**04**

**PART FOUR**

# **Test Programs**



# 04 / Test Programs

- Four public test programs, respectively named mp2\_N with N = 1,... , 4 , source code at /user/mp2\_N.c

1. Run the command to launch all tests at once, and the output will be saved to mp2\_N.out files.

```
$ ./run_mp2.py
...
== Test mp2_1 == (1.3s)
== Test mp2_2 == (0.4s)
== Test mp2_3 == (1.0s)
== Test mp2_4 == (1.0s)
```

2. To run an individual test program instead, run “make qemu” to enter the xv6 shell and “run mp2\_N”.

```
$ make clean
$ make qemu
...

xv6 kernel is booting

hart 2 starting
hart 1 starting
init: starting sh
$ mp2_1
```



**05**

**PART FIVE**

**Submission**

# 05 / Submission

## **MP2 assignment deadline : April 4, 23:59:00**

### **1. Source code:**

Submit your d08922025.zip to “Machine Problem 2”

### **2. Report:**

Submit one PDF file to "Report" named d08922025\_mp2\_report.pdf, for example.

### **3. Bonus report (optional):**

Submit one PDF file to "Bonus Report" named d08922025\_mp2\_bonus.pdf, for example.